

Machine Learning with TensorFlow & Keras

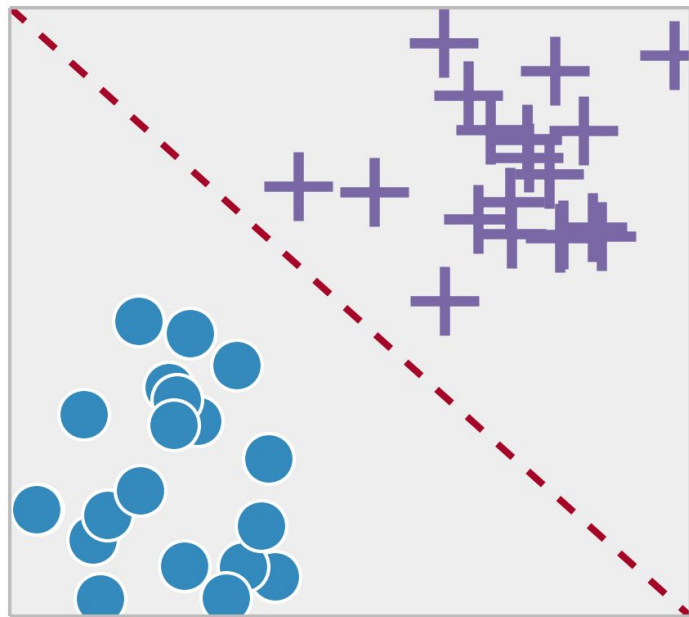
Ganesh Katrapati - PhD Candidate @ IIITH, EC Member - Swecha

ML ?

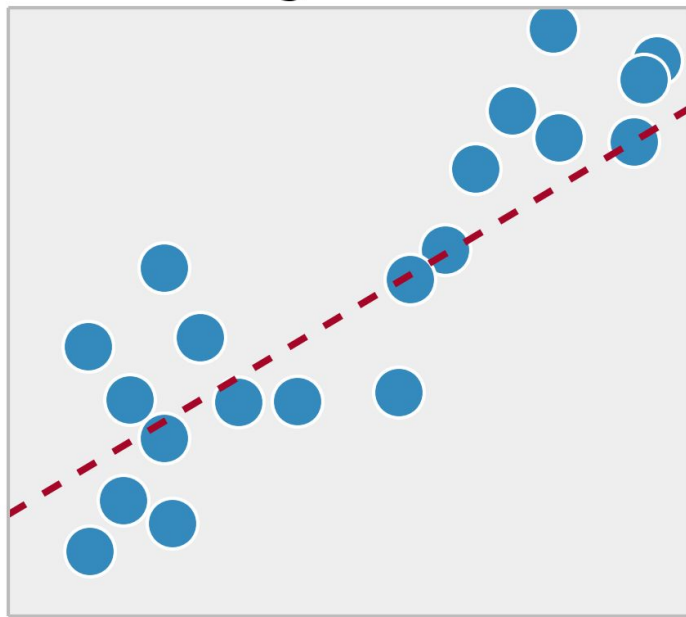
Look at the past to predict **a possible** future.

Two Problem Classes (mostly)

Classification



Regression



Regression

- Age vs Height - I have data for ages 1-20, 40-60. Can I form a pattern for 20-40 ?
- Experience vs Salary
- Global Temperature vs FF Emission.
- Price Rise vs Demand.
- Change vs Risk (degree)

Classification

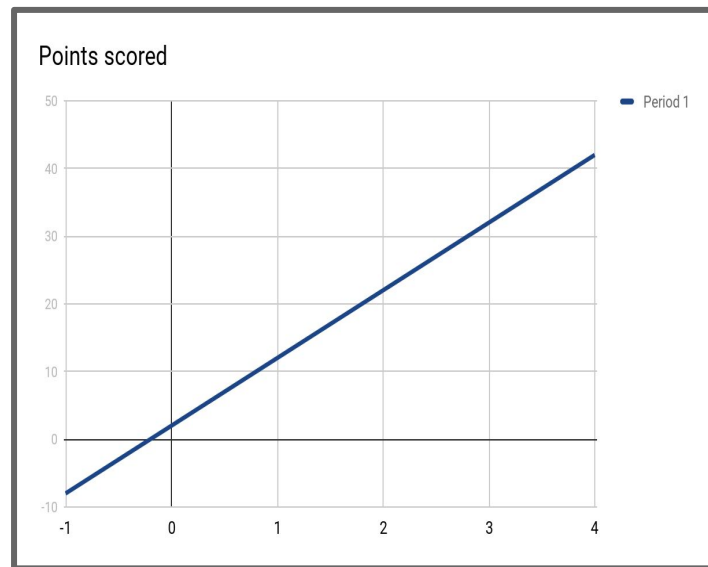
- Apples vs Oranges : I have a *labeled dataset* for Apples vs Oranges. Given a fruit, can I decide ?
- Risk vs Not-A-Risk
- Carcinogenic or Not
- Disease Classification

The General Idea

- So, in general, both Classification and Regression are about learning a **line**.
- Lets go back to math. What is a line ?

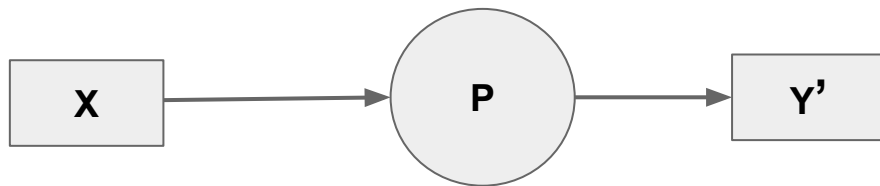
$$Y = MX + C$$

- We know : X (input), Y(output)
- We don't know : M & C (Slope & Intercept)
- Lets call these : **weights**



The Machine

$$\mathbf{Y} = \mathbf{M}\mathbf{X} + \mathbf{C}$$



- Algorithm
 - For every input **X**, we run it through the Perceptron (**P**)
 - **P** outputs some output **Y'**
 - Our goal is to make **Y'** as similar as possible to **Y**.

The Algorithm



- The central goal : minimise loss function (J)
 - $J = Y - Y'$
- What is learning ?
 - Learn the required weights which minimise J.
 - For each turn, adjust weights proportional to the loss encountered.
 - Usually, learning is not one-shot. So, learning is according to a ***learning rate***.

Lets Complicate it a little bit :)

- Input can have more than one property
 - For apples and oranges, it can be color & shape for example
 - $Y = W_1X_1 + W_2X_2 + C$

- Consider a case where
Drawing a line is useless.

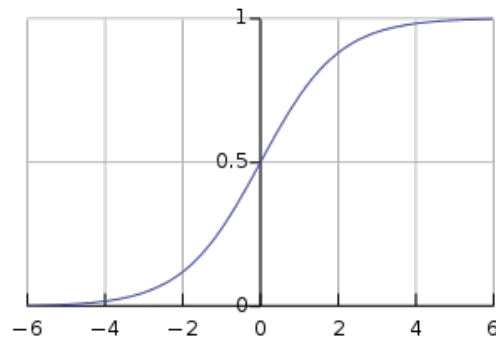
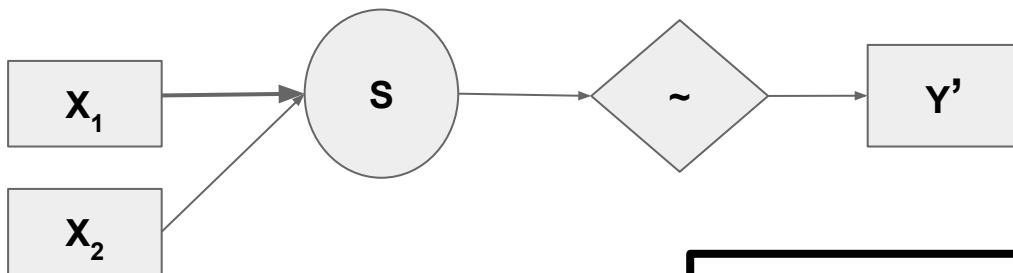
Linear vs. nonlinear problems



A Super Power : Non-Linearity

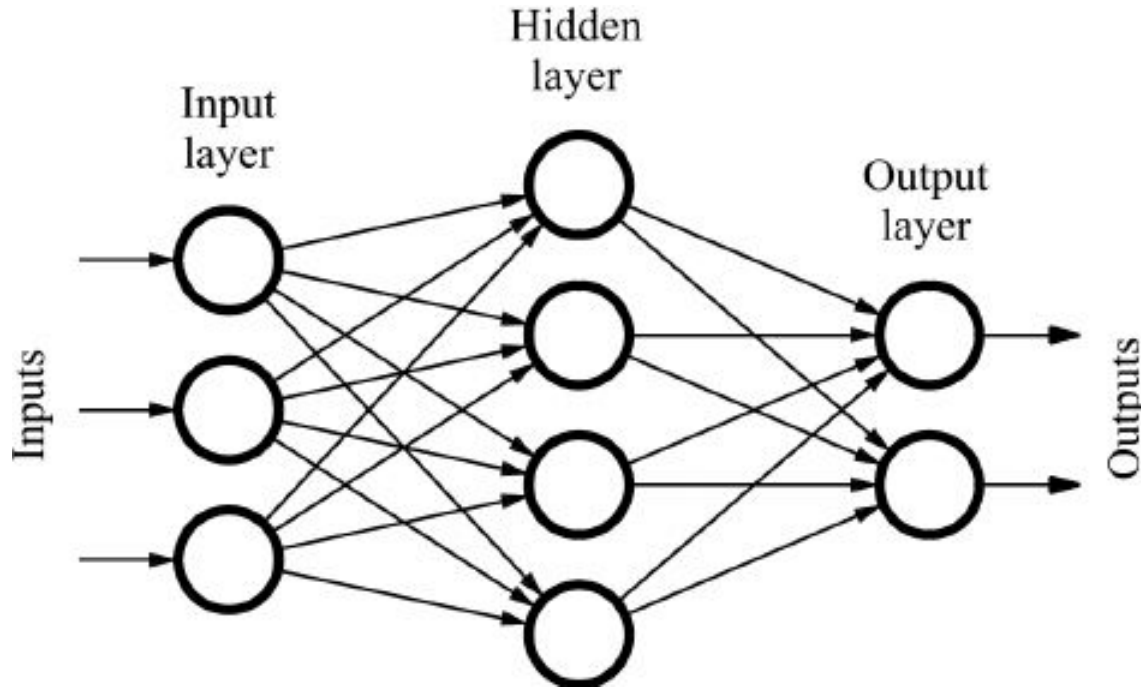
- We have to send our output (which is a sum) through a non-linear function (eg: sigmoid)

$$Y = F^{\sim}(W_1X_1 + W_2X_2 + C)$$



ML is a generic function approximation algorithm

Put it all together = NN



- At each 'neuron', we determine its contribution to the loss function.
- When **units** are stacked in the shown manner (**layers**). Complex functions can be learnt
- When we do this for all the units, we approach a desired state
- This is called **Backpropagation**.

Lets get started

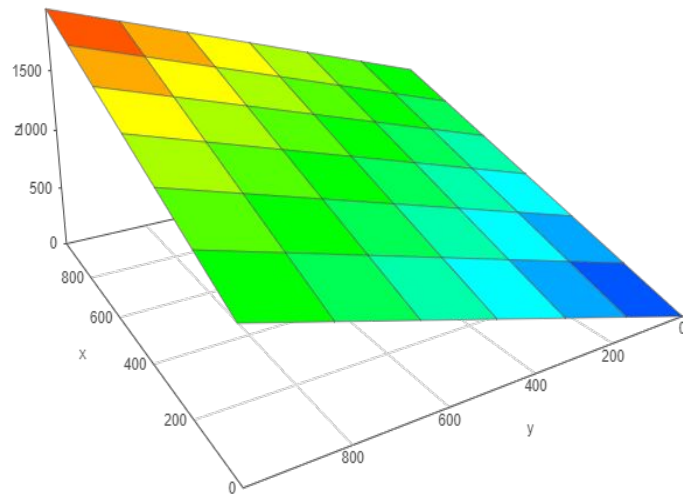
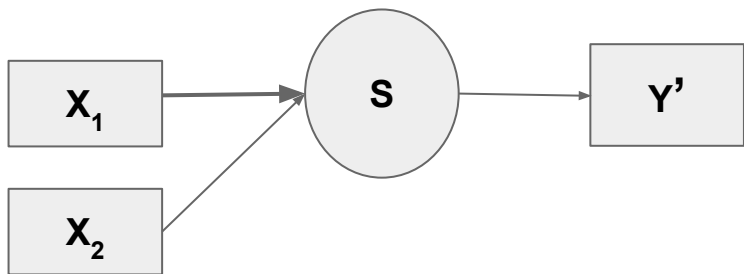
- TensorFlow -> A library which is very customisable to the extent of individual layers and units to learn and use neural nets
- **Keras** : TensorFlow for humans (kind-of).

Setting Up

- Python
- `pip install tensorflow`
- `pip install keras`
- Windows ?
 - Install Anaconda 3.6
 - `Conda install mingw libpython`
 - `Conda install tensorflow`
 - `Pip install keras`

Our First Neural Net

- Problem : Let us make our NN *learn addition*.
- Inputs : X_1, X_2
- Output : Y (where $Y = X_1 + X_2$)
- The data will be linearly separable.



Anatomy of a Keras Script

Imports

```
from keras.models import Sequential  
from keras.layers import Dense, Activation
```

The basic keras api.

Dense = Summing unit

Activation = A function upon dense.

Create a model

```
model = Sequential()  
model.add(Dense(1, input_dim=2))
```

Input Dimensions = 2 (X1, X2)

Output Dimensions = 1 (Y)

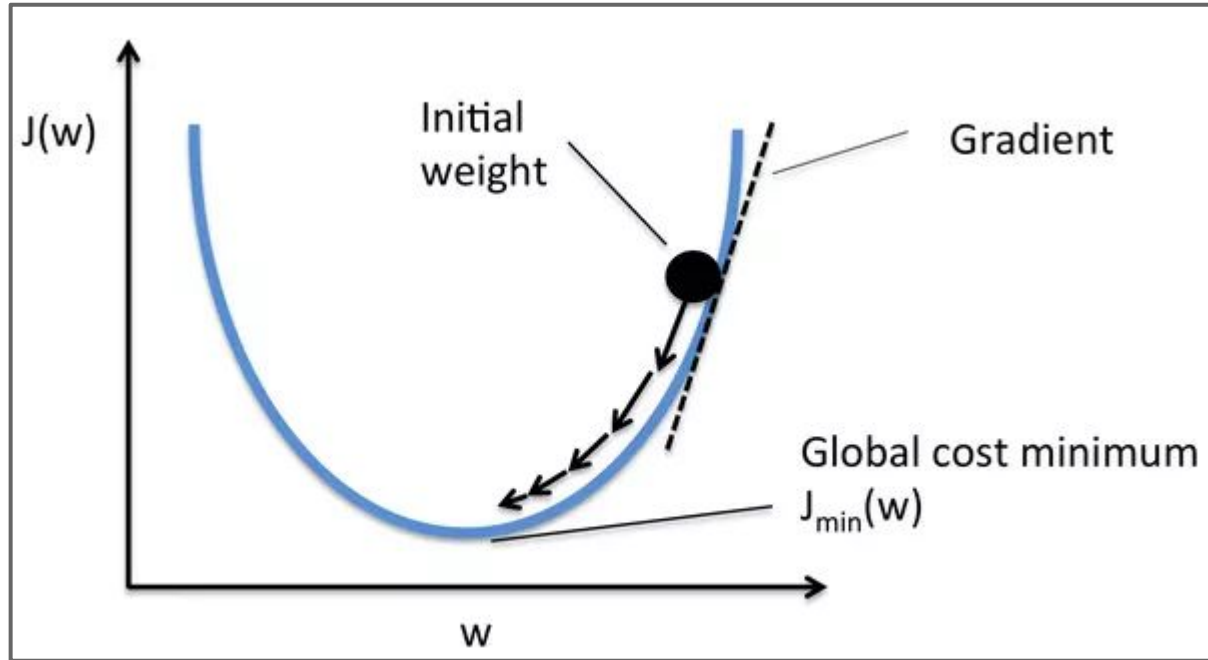
Compile the model

```
model.compile(optimizer='sgd',  
              loss='mse')
```

SGD = Stochastic Gradient Descent

MSE = Mean Squared Error

A Side Note : SGD & MSE



- Error (J) will be better if it is squared $(Y - Y')^2$ because it will be smoother.
- Gradient Descent is a process of slowly crawling down to the lowest J possible.

Anatomy of a Keras Script - Training

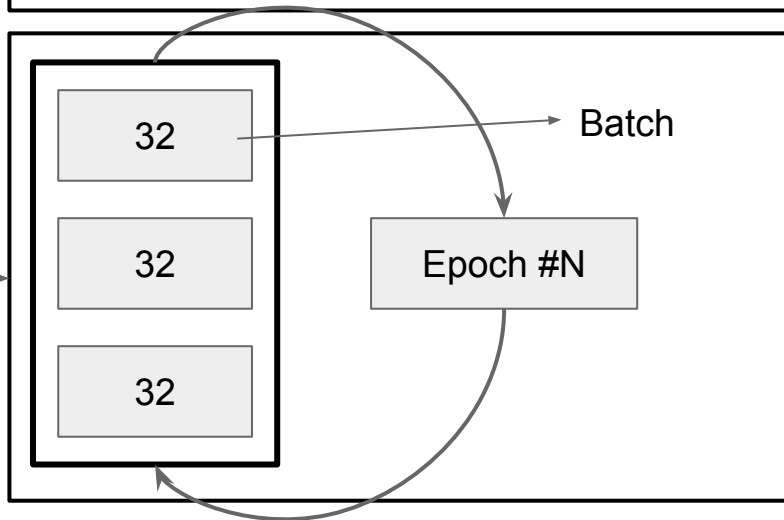
Generate Training data

```
import numpy as np  
TRAIN_SIZE = 1000  
X = np.random.random((TRAIN_SIZE, 2))  
Y = X.sum(axis = 1)
```

Import numpy
Generate X (1000x2) matrix
Sum columns to get Y ($X_1 + X_2$)

Train

```
model.fit(X, Y, epochs=10, batch_size=32)
```



Anatomy of a Keras Script - Testing

Generate Training data

```
TEST_SIZE = 100  
X_test = np.random.random((TRAIN_SIZE,  
2))  
Y_test = X_test.sum(axis = 1)
```

Evaluate(Optional)

```
score = model.evaluate(X_test, Y_test)  
print("Error is", score * 100, "%")
```

Score is the Mean Squared Error over all the samples

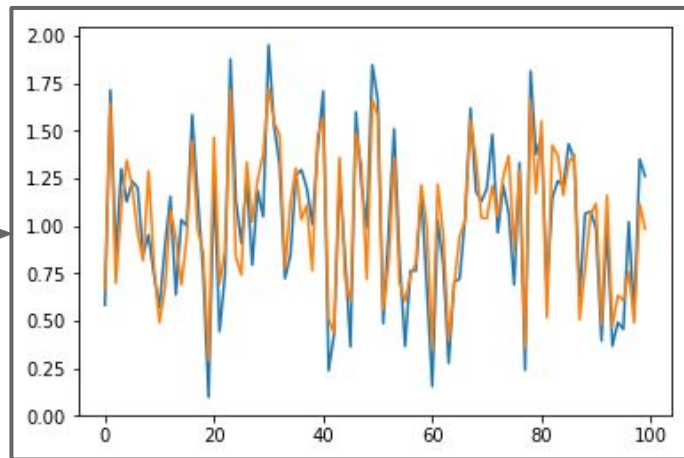
Get the predicted Y's

```
Y_pred = model.predict(X_test)  
print(Y_test, Y_pred)
```

Bonus : Plot

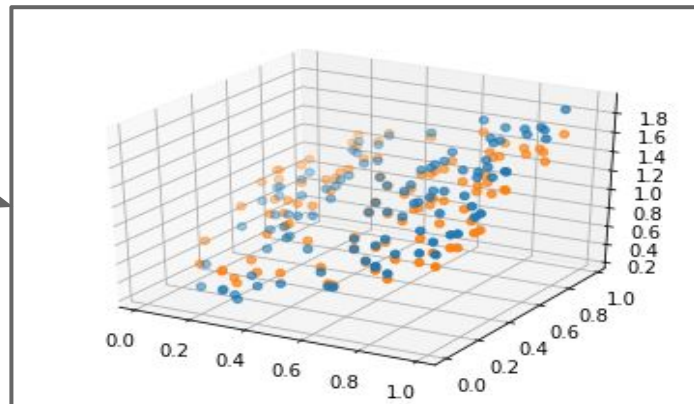
Y_test vs Y_pred

```
import matplotlib.pyplot as plt  
plt.plot(range(0,len(Y_test)), Y_test)  
plt.plot(range(0,len(Y_pred)), Y_pred)  
plt.show()
```



Get the predicted Y's

```
from mpl_toolkits.mplot3d import Axes3D  
fig = plt.figure()  
ax = fig.add_subplot(111, projection='3d')  
ax.scatter(X_test[:,0], X_test[:,1], Y_test)  
ax.scatter(X_test[:,0], X_test[:,1], Y_pred)  
plt.show()
```

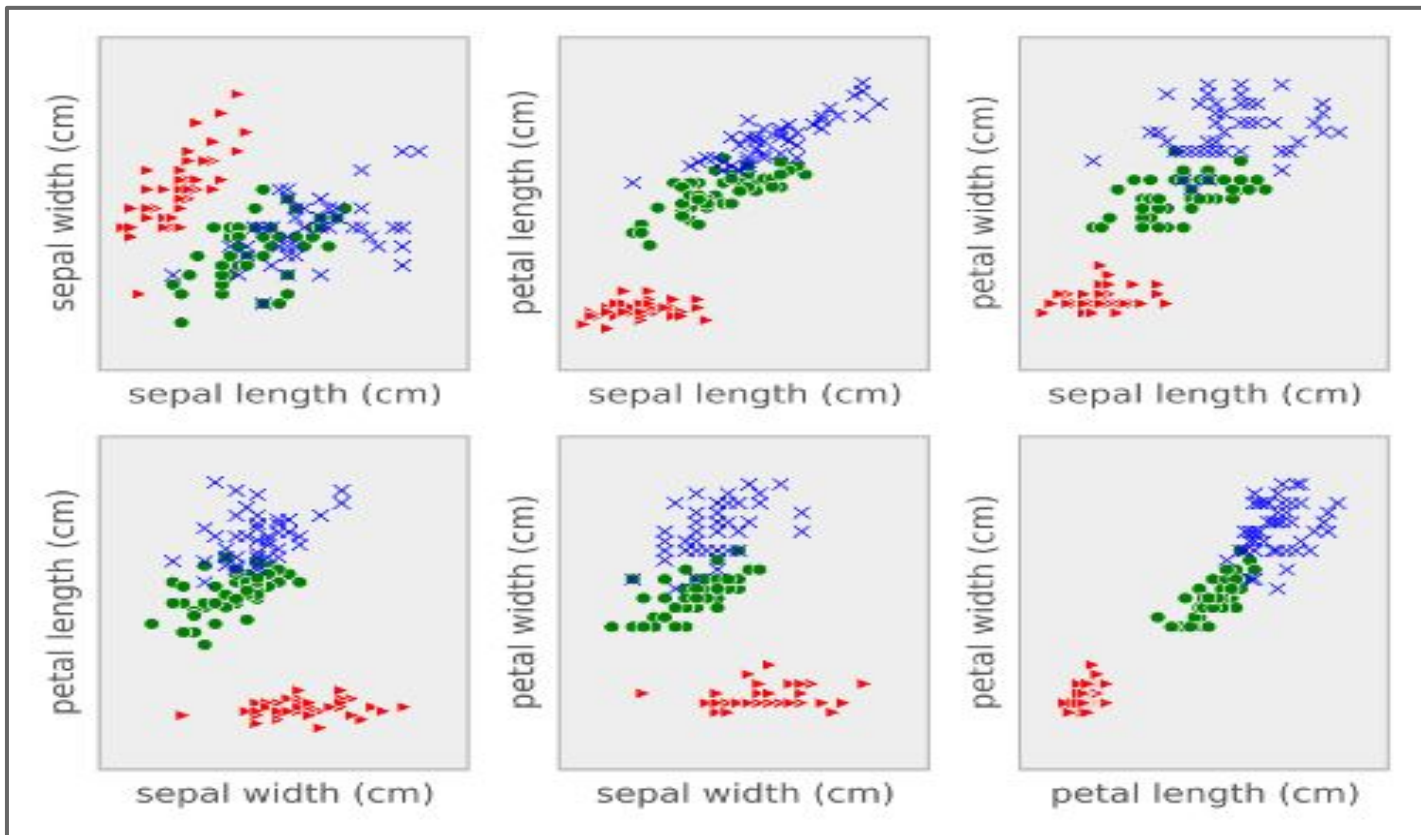


Kudos !. You have built the
great-great-great-great-----
-grandmother of Skynet.


Take 2 : Multi-Class Classification

- Problem : Classify IRIS flowers into subtypes.
- **Inputs** : Sepal Length, Sepal Width, Petal Length, Petal Width
- **Output : Class Name**
 - Iris Setosa
 - Iris Versicolor
 - Iris Virginica

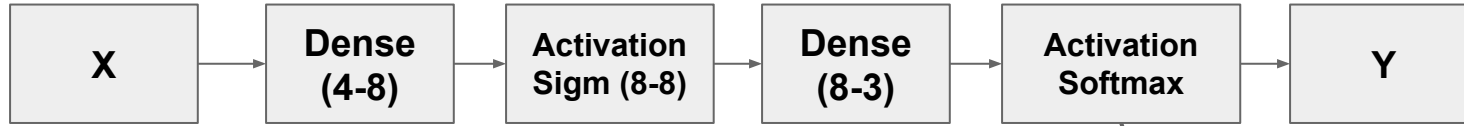
IRIS



IRIS - Network Design



IRIS - Network Design



New Model

```
model = Sequential()  
model.add(Dense(units = 8, input_dim=4))  
model.add(Activation('sigmoid'))  
model.add(Dense(units = 3))  
model.add(Activation('softmax'))
```

Softmax function gives relative likelihood of which class a sample belongs to. For example : X might be (75%, 15%, 10%) if Iris (Setosa, Virginica, Versicolor)

Why More Layers ?

- For more complex problems (data is not simply separable), more layers allow learning these complex decision boundaries
- Any NN with $\#L > 2$ is called a Deep Neural Net. aka “Deep Learning”
- Usually, each layer learns a level of generalisation.
 - Say for face recognition,
 - Layer 1 learns edges & lines
 - Layer 2 learns facial components
 - Layer 3 learns to distinguish faces themselves.